

4 Stage Pipelined Multimedia Unit Design with VHDL

ESE 345

Fall 2024

Professor Mikhail Dorojevets

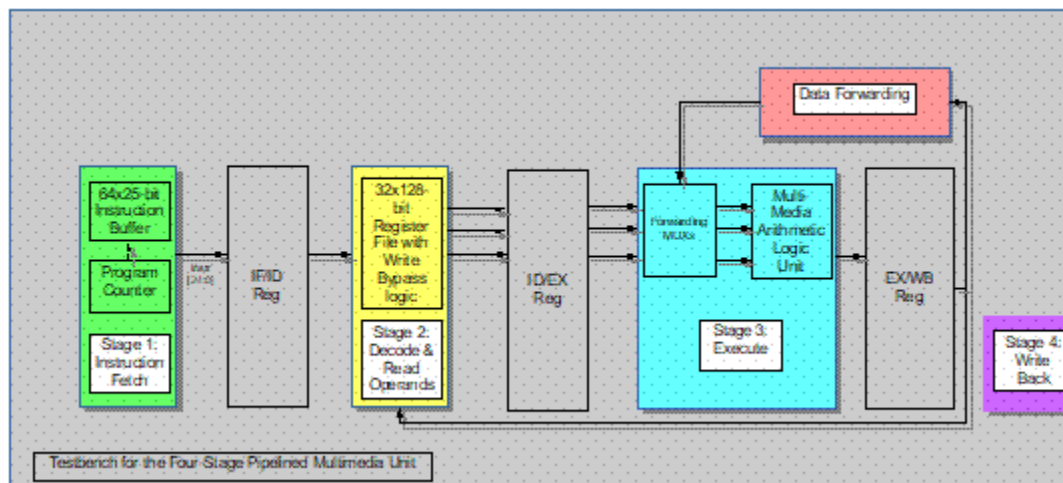
Naafiul Hossain and Muhammed Sharjeel

Introduction and Goal:

ESE 345, *Computer Architecture*, is a course at Stony Brook University that focuses on processor design, memory structure, and advanced system design concepts. The course culminates in a capstone project: designing a pipelined multimedia processing unit using a hardware description language of our choice—in this case VHDL. Additionally, we were tasked with developing a custom assembler program tailored to our hardware design—C++.

The final evaluation required us to test, validate, and present our work to a teaching assistant. During the presentation, we demonstrated our assembler's functionality by converting a MIPS-style assembly program—provided by the TA—into binary machine code. We showcased the execution of the code and highlighted the memory addresses where results were stored.

The project is broken into eight key units as shown in this illustration:



Multimedia ALU: Takes up to three inputs from the Register File, and calculates the result based on the current instruction to be performed. The ALU must be implemented as behavioral model in VHDL or continuous assignment (dataflow models in Verilog).

Register File: The register file has 32 128-bit registers. On any cycle, there can be 3 reads and 1 write. When executing instructions, each cycle two/three 128-bit register values are read, and one 128-bit result can be written if a write signal is valid. This register write signal must be explicitly declared so it can be checked during simulation and demonstration of your design. The register module must be implemented as a behavioral model in VHDL (dataflow/RTL model in Verilog).

Instruction Buffer: The instruction buffer can store 64 25-bit instructions. The contents of the buffer should be loaded by the testbench instructions from a test file at the start of simulation. On each cycle, the instruction specified by the Program Counter (PC) is fetched, and the value of PC

is incremented by 1. The Instruction Buffer module must be implemented as a behavioral model in VHDL (dataflow/RTL model in Verilog).

Forwarding Unit: Every instruction must use the most recent value of a register, even if this value has not yet been written to the Register File. Be mindful of the ordering of instructions; the most recent value should be used, in the event of two consecutive writes to a register, followed by a read from that same register. Your processor should never stall in the event of hazards. Take extra care of which instructions require forwarding, and which ones do not. Namely, NOP and the instructions with Immediate fields do not contain one/two register sources. Only valid data and source/destination registers should be considered for forwarding.

Four-Stage Pipelined Multimedia Unit: Clock edge-sensitive pipeline registers separate the IF, ID, EXE, and WB stages. Data should be written to the Register File after the WB Stage. All instructions (including li) take four cycles to complete. This pipeline must be implemented as a structural model with modules for each corresponding pipeline stages and their interstage registers. Four instructions can be at different stages of the pipeline at every cycle.

Testbench: This module loads the instruction buffer using data loaded from a file, begins simulation, and upon completion, compares the contents of the register file to a file containing the expected results. This expected results file does not need to be auto-generated. Instead, this can be manually entered when designing a test program. This must be implemented as a behavioral model.

Assembler: This is a separate program written in any language your team prefers (i.e. Java, C++, Python). Its purpose is to convert an assembly file to the binary format for the Instruction Buffer. This assembler does not need to be robust, and can assume very specific syntax rules that you as a team decide.

Results File: This file must show the status of the pipeline for each cycle during program execution. It should include the opcodes, input operand, and results of the execution of instructions, as well as all relevant control signals and forwarding information. This should be carried out by your testbench.




Instruction Format

The instruction format for the Multimedia Unit is designed to support three distinct types of instructions: I-type, R3-type, and R4-type, each with a unique structure. I-type instructions are 24 bits long and are identified by a 1 in the 23rd bit. They include a 16-bit immediate value (bits 20–5) that is loaded into the

rd register specified by bits 4–0. Additionally, bits 22–21 indicate the byte of the rd register where the immediate value will be placed. This format allows for efficient handling of immediate data operations.

R3-type instructions, also 24 bits, are characterized by the 00 in bits 23–22. These instructions include a 7-bit opcode (bits 21–15) that specifies the operation to be performed, and three registers: rs2, rs1, and rd, represented by bits 14–10, 9–5, and 4–0, respectively. This format supports operations that require two source registers and one destination register, making it suitable for arithmetic and logical computations. The clear structure ensures accurate decoding and execution within the pipeline.

R4-type instructions expand on the R3 format by adding a field for additional functionality. Identified by 01 in bits 23–22, this format includes MA/MS/I/h bits in positions 21–20 to select specific operations. It uses four registers: rs3 (bits 19–15), rs2 (bits 14–10), rs1 (bits 9–5), and rd (bits 4–0), enabling more complex instructions. The structured approach across these three formats allows the pipeline to efficiently process diverse instruction sets, optimizing performance and reducing execution time.



Logic Implementation:


The system implementation involved several critical components working in harmony to create a functional Pipelined Multimedia Unit (PMU). The testbench played a central role, programming the instruction buffer, providing clock and reset signals, and validating results by comparing the processor's memory with expected outputs.

It operated in three states (resetting, program control, and normal operation) feeding instructions into program memory during the programming state and generating debug messages for any discrepancies. The instruction buffer managed up to 64 instructions, propagating instruction flags and signaling when the program was complete. It ensured proper sequencing of instructions through pipelining stages, using a forwarding register for spacing and accurate handoff of instruction data.

The register file handled data storage and retrieval for arithmetic operations while mitigating hazards through write-back checks. It outputs data based on instruction types, supporting seamless integration with the pipeline. Forwarding registers in the instruction fetch and decode stages ensured proper data flow, maintaining pipeline spacing and forwarding valid instruction flags. Throughout, verification strategies emphasized transparency and accuracy, leveraging step-by-step memory checks against expected results. This robust design, paired with effective debugging and validation, ensured the PMU met its functional and performance objective

Assembler:

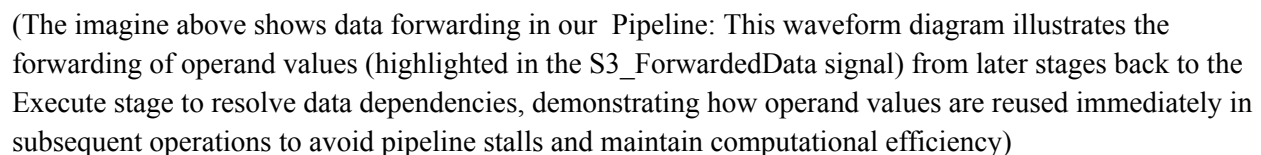
Our program, written in C++, is a MIPS assembly to machine code translator. It reads MIPS assembly instructions from an input file (assemblerN.txt), converts each instruction or register to its corresponding machine code using a predefined mapping stored in a std::map, and writes the resulting machine code to an output file (Nmachine_code.txt). The code includes a function to translate decimal numbers into 16-bit binary for immediate values and another function to map assembly instructions or registers to machine code strings. Each line in the input file is processed word by word: numeric values are converted to binary, and assembly instructions or registers are matched in the map for translation. The resulting machine code for each line is padded to 25 characters before being written to the output file. The program uses file handling, string manipulation, and the STL (Standard Template Library) to achieve efficient translation.



Forwarding:


The forwarding logic in the execute module is implemented by dynamically comparing specific fields of the current instruction's opcode (EXin_opcode) with the forwarded instruction's opcode (opcode_fwd). It identifies data hazards by matching relevant bits and uses the forwarded result (rd_fwd) to supply the appropriate operand (rs1, rs2, or rs3). The logic categorizes hazards into three types: **LI hazards**, which focus on the lower bits of the opcode to detect dependencies on rs1; **R4 hazards**, which resolve dependencies for multi-operand instructions by comparing higher-order opcode bits with all three operands; and **R3 hazards**, which handle simpler two-operand instructions by selectively forwarding to rs1 or rs2. If a match is found and the forwarded instruction is valid (Valid_Instruction_In_FWD = '1'),

This logic works because it resolves data dependencies dynamically during the execution stage of a pipeline, eliminating the need for stalls. By forwarding results directly from a later stage to an earlier stage, the module avoids delays caused by waiting for register writes. This implementation is efficient and ensures the pipeline continues executing instructions without interruption. The decision-making process is driven by specific opcode patterns, which precisely identify when and where forwarding is needed. This mechanism works because it leverages the structure of pipelined execution, where instructions in different stages have predictable dependencies, allowing forwarding to mitigate data hazards without additional complexity.



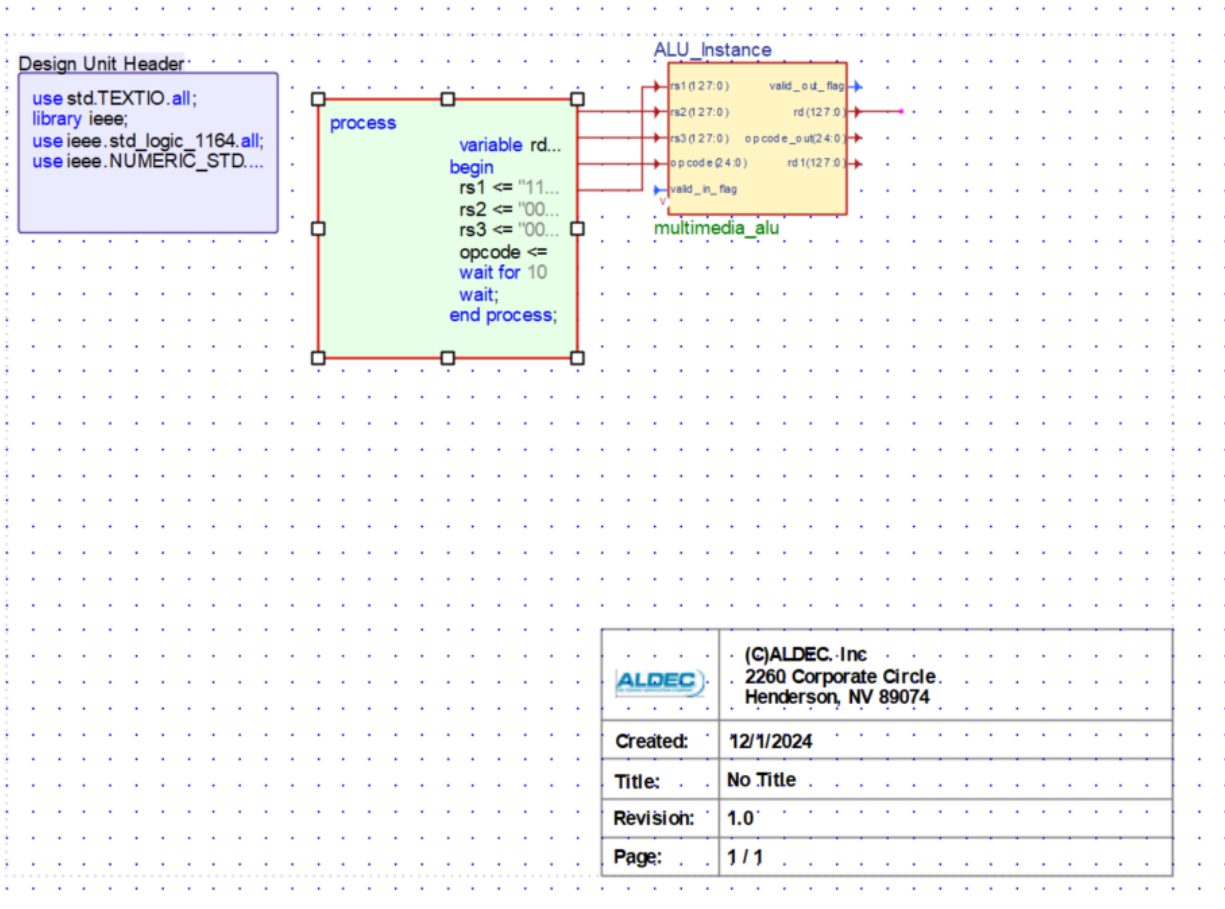
In this waveform that we generated, we're looking at a snapshot of a microprocessor's pipelining process, which allows multiple instructions to be processed simultaneously at different stages. Here, pipelining is illustrated through the S1_PC and S1_NextInstruction signals, showing where in the pipeline instructions are currently positioned. The S1_PC indicates the program counter for the instruction being processed, and S1_NextInstruction represents the address of the next instruction to be fetched. Meanwhile, the S2_Instruction seems to be idle at this point, indicating that it's either waiting for new instructions to

arrive from the fetch stage or there is a stall in the pipeline. This simultaneous handling of instructions at different stages optimizes throughput and efficiency, as while one instruction is being decoded, another can be fetched, reducing idle times and increasing the overall speed of the processing.

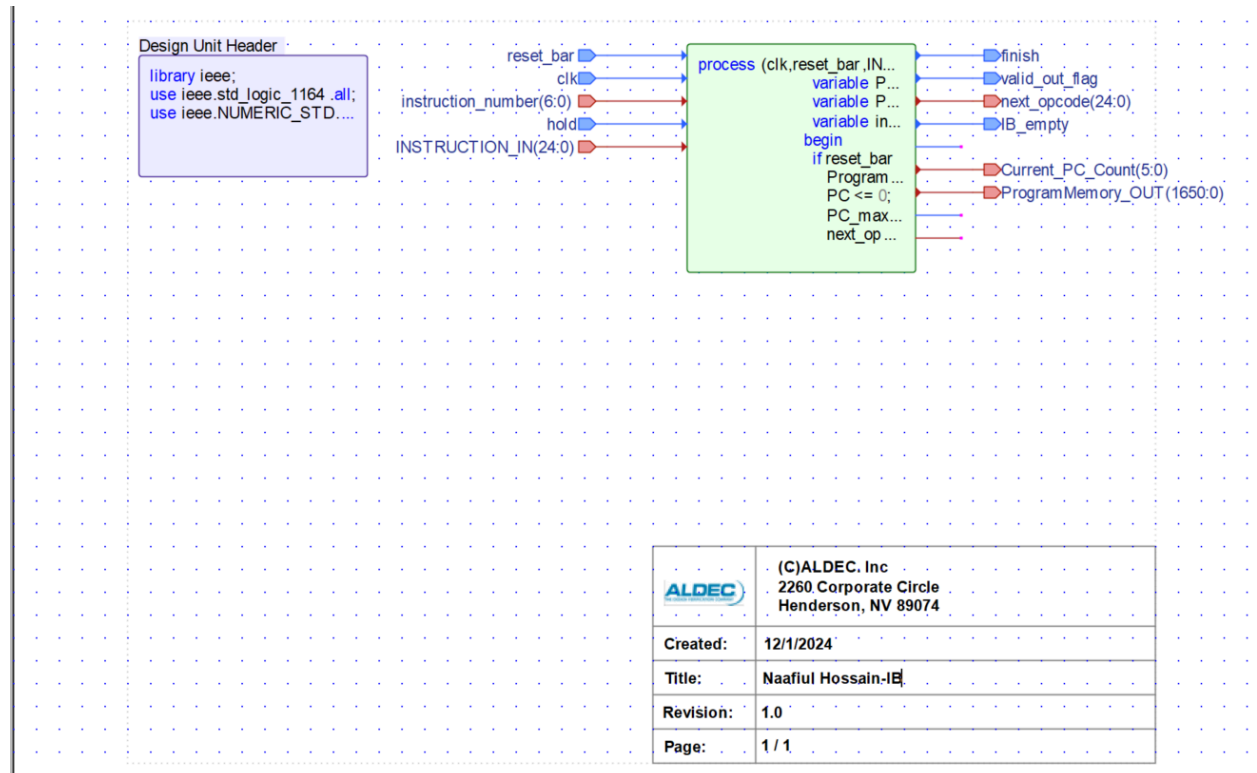


Block Diagrams

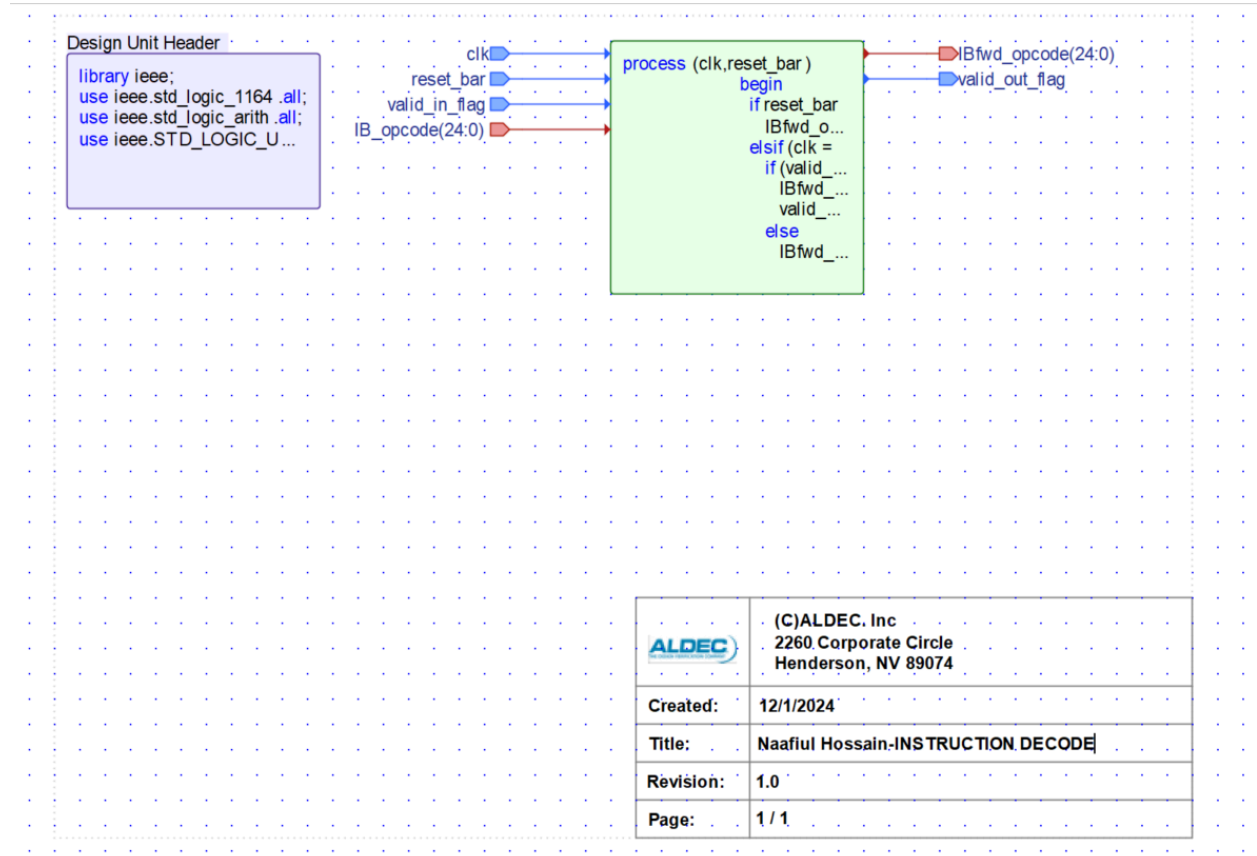
Alu:



Instruction Buffer:



Instruction Decode:



Design Unit Header

```

library ieee;
use ieee.std_logic_1164.all;
use ieee.std_logic_arith.all;
use ieee.STD_LOGIC_U...


```

```

process (EXin_rs1, EXin_r...
    variable fo...
    variable fo...
begin
    if reset_bar
    else
        if (valid ...
            if not (...
                if EX...
                    for ...

```

	(C)ALDEC, Inc 2260 Corporate Circle Henderson, NV 89074
	Created: 12/1/2024
	Title: Naafli Hossain-EXECUTE
	Revision: 1.0
	Page: 1 / 1



Design Unit Header

```
library ieee;
use ieee.std_logic_1164.all;
use ieee.NUMERIC_STD ...
```

Architecture Declaration

```
--Added by Active-HDL. D...
type RF is array(0 to 31) of
--End of extra code.
```

```
clk
opcode(24:0)
wb_opcdata(24:0)
reset_bar
register_write
wb_reg_data(127:0)
valid_in_flag
valid_in_flag_WB
IB_empty
hold
```

```
process (clk, reset_bar, re...
begin
if (opcode(
reg_sel_A
else
reg_sel_A
reg_sel_B
reg_sel_C
end if;
if reset_bar
reg_A <=
reg_B <=
reg_C <=
reg <= (...
reg(0) <...
end if;
if (IB_empty
reg_A <=
reg_B <=
reg_C <=
opcode_...
valid_ou...
else
if (valid_...
valid_...
if opco...
reg_A
elsif o...
reg_A
reg_B
reg_C
elsif o...
reg_8_out(127:0)
reg_5_out(127:0)
reg_6_out(127:0)
reg_4_out(127:0)
r...
r...
when
r...
when
end ...
case
when
r...
r...
when
end ...
end if;
end if;
```

```
reg_3_out(127:0)
reg_19_out(127:0)
reg_27_out(127:0)
reg_29_out(127:0)
reg_31_out(127:0)
reg_out_all(4095:0)
reg_30_out(127:0)
reg_28_out(127:0)
reg_23_out(127:0)
reg_25_out(127:0)
reg_26_out(127:0)
reg_24_out(127:0)
reg_21_out(127:0)
reg_22_out(127:0)
reg_20_out(127:0)
reg_11_out(127:0)
reg_15_out(127:0)
reg_17_out(127:0)
reg_18_out(127:0)
reg_16_out(127:0)
reg_13_out(127:0)
reg_14_out(127:0)
reg_12_out(127:0)
reg_7_out(127:0)
reg_9_out(127:0)
reg_10_out(127:0)
reg_8_out(127:0)
reg_5_out(127:0)
reg_6_out(127:0)
reg_4_out(127:0)
reg_A(127:0)
valid_out_flag
reg_1_out(127:0)
reg_2_out(127:0)
reg_0_out(127:0)
reg_C(127:0)
opcode_out(24:0)
reg_B(127:0)
```



(C)ALDEC, Inc.
2260 Corporate Circle
Henderson, NV 89074

Created:

12/1/2024

Title:


NAAFIUL HOSSAIN RH

Revision:


1.0

Waveform Stimulations:

Load:



NOP:



SLHI: (ex. shift 6)

AND:

FFFF0001F1F1F1747474748989898989
FFFF0000FFFF0000FFFF0000FFFF0000
00000000000000000000000000000000
1828000
FFFF0000F1F100007474000089890000

BCW:

FFFF0001F1F1F1747474748989898989
00000000000000000000000000000000
00000000000000000000000000000000
1830000
89898989898989898989898989898989

MAXWS:

FFFF0001F1F1F1747474748989898989
0200000600000000400000100100007FD
00000000000000000000000000000000
1838000
0200000600000000474747489100007FD

MINWS:

FFFF0001F1F1F1747474748989898989
0200000600000000400000100100007FD
00000000000000000000000000000000
1838000
0200000600000000474747489100007FD

MLHU:

00FF0F00FFDC001409C4003FC0618142
F8F8F8FFC7C7C7FE3E3E3FF1F1F1F1
00000000000000000000000000000000
1848000
0E96F100000F9FD8000FBC4F7A28D122

MLHCU:

00FF0F00FFDC001409C4003FC0618142
F8F8F8FFC7C7C7FE3E3E3FF1F1F1F1
F8F8F8FFC7C7C7FE3E3E3FF1F1F1F1
1857400
0001B3000000024400000723000EA47A

OR:

FFFF0001F1F1F1747474748989898989
FFFF0000FFFF0000FFFF0000FFFF0000
00000000000000000000000000000000
1858000
FFFF0001FFFFFF174FFFF7489FFFF8989

CLZH:

00FF0F00FFDC001409C4003FC0618142
F8F8F8FFC7C7C7FE3E3E3FF1F1F1F1
F8F8F8FFC7C7C7FE3E3E3FF1F1F1F1
1860000
0008000400000000B0004000A00000000

RLH:

00FF0F00FFDC001409C4003FC0618142
00000001000200060009000B000D000F
F8F8F8FFC7C7C7FE3E3E3FF1F1F1F1
1868000
00FF1E00FF7305008813F801380C40A1

SFWU:

FFFF0001F1F1F1747474748989898989
0000FFFF0F0F103F45450000FFFF0387
00000000000000000000000000000000
1870000
0001FFFE1D1D1ECBD0D08B77767579FE

SFHS:

FFFF0001F1F1F1747474748989898989
0000FFFF0F0F103F45450000FFFF0387
00000000000000000000000000000000
1878000
0001FFFE1D1E1ECBD0D18B77767679FE

4.2 Multiply-Add and Multiply-Subtract R4-Instruction Format

Signed Integer Multiply-Add Low with Saturation:

00000004000000004800000007FFFFFFD
0000000600000000400000001000007FD
000000080000000010000800000006FFE
1000000
00000034000000008800000007FFFFFFF

Signed Integer Multiply-Add High with Saturation:

00000004000000004800000007FFFFFFD
000600000000400000001000007FD0000
00080000000010000800000006FFE0000
1100000
00000034000000008800000007FFFFFFF

Signed Integer Multiply-Subtract Low with Saturation:

0000000480000000800000007FFFFFFD
0000000600000000400000001000007FD
000000080000000010000800000006FFE
1200000
800000007FFFFFFC800080007C815FF7

Signed Integer Multiply-Subtract High with Saturation:

0000000480000000800000007FFFFFFD
000600000000400000001000007FD0000
00080000000010000800000006FFE0000
1300000
800000007FFFFFFC800080007C815FF7

Signed Long Multiply-Add Low with Saturation:

75FFFEFDFF53F4FF08000000000400FEF
00000072000000030000000083289192
9700000100000001000000000832910
1400000
75FFFEFDFF53F4FF38000000000000000

Signed Long Multiply-Add High with Saturation:

75FFFEFDFF53F4FF08000000000400FEF
00000003000000728328919200000000
00000001970000010083291000000000
1500000
75FFFEFDFF53F4FF38000000000000000

Signed Long Multiply-Subtract Low with Saturation:

800000007FFFFFFF0000000083289192
000000007FFFFFFF0000000083289192
000000007F000000F000000000832910
1600000
407FFFF97F000000E003FF64453991672

Signed Long Multiply-Subtract High with Saturation:

75FFEFDFF53F4FF08000000000400FEF
0000007200000000300000000083289192
97000001000000001000000000832910
1700000
75FFF00EB73F4F7E8000000000400FEF

Conclusions/Final statements :

In conclusion, our multimedia processing unit effectively processes instructions, adhering to the principles of pipelining and using techniques such as forwarding when facing hazards. Utilizing VHDL allowed us to implement complex logic for handling various operational scenarios.

While most instructions were executed efficiently, some required additional verification to ensure accuracy and performance optimization. This project not only reinforced our understanding of complex architecture concepts but also provided practical experience in troubleshooting and enhancing processor functionality with VHDL.

Through this course and final project, we have gained significant insights into the operational challenges and considerations involved in designing and implementing a pipelined system, laying a solid foundation for future explorations in computer architecture.